

Bike Sharing in Washington DC

August 11, 2020

1 Bike Sharing in Washington DC Data Analysis

Link to the data: <https://www.kaggle.com/marklvl/bike-sharing-dataset>

Data Collection Date: 2011 and 2012

Data Source: 'Capital Bikeshare' <https://www.capitalbikeshare.com/system-data>

Possible Questions to explore:

1. How do variables like time of day, season, and weekday/weekend affect the number of bike rentals in a day?

```
[81]: import numpy as np
import pandas as pd
import statsmodels.api as sm
```

```
[82]: df = pd.read_csv('hour.csv')
df = df.rename(columns={'cnt': 'count'})
df['workingday'] = df['workingday'].replace({0: False, 1: True})
```

```
[83]: df.head()
```

```
[83]:
```

| | instant | dteday | season | yr | mnth | hr | holiday | weekday | workingday | \ |
|---|---------|------------|--------|----|------|----|---------|---------|------------|---|
| 0 | 1 | 2011-01-01 | 1 | 0 | 1 | 0 | 0 | 6 | False | |
| 1 | 2 | 2011-01-01 | 1 | 0 | 1 | 1 | 0 | 6 | False | |
| 2 | 3 | 2011-01-01 | 1 | 0 | 1 | 2 | 0 | 6 | False | |
| 3 | 4 | 2011-01-01 | 1 | 0 | 1 | 3 | 0 | 6 | False | |
| 4 | 5 | 2011-01-01 | 1 | 0 | 1 | 4 | 0 | 6 | False | |

| | weathersit | temp | atemp | hum | windspeed | casual | registered | count |
|---|------------|------|--------|------|-----------|--------|------------|-------|
| 0 | 1 | 0.24 | 0.2879 | 0.81 | 0.0 | 3 | 13 | 16 |
| 1 | 1 | 0.22 | 0.2727 | 0.80 | 0.0 | 8 | 32 | 40 |
| 2 | 1 | 0.22 | 0.2727 | 0.80 | 0.0 | 5 | 27 | 32 |
| 3 | 1 | 0.24 | 0.2879 | 0.75 | 0.0 | 3 | 10 | 13 |
| 4 | 1 | 0.24 | 0.2879 | 0.75 | 0.0 | 0 | 1 | 1 |

Description of the data can be found here: <https://www.kaggle.com/marklvl/bike-sharing-dataset>

```
[84]: df.describe()
```

```
[84]:
```

| | instant | season | yr | mnth | hr \ |
|-------|------------|--------------|--------------|--------------|--------------|
| count | 17379.0000 | 17379.000000 | 17379.000000 | 17379.000000 | 17379.000000 |
| mean | 8690.0000 | 2.501640 | 0.502561 | 6.537775 | 11.546752 |
| std | 5017.0295 | 1.106918 | 0.500008 | 3.438776 | 6.914405 |
| min | 1.0000 | 1.000000 | 0.000000 | 1.000000 | 0.000000 |
| 25% | 4345.5000 | 2.000000 | 0.000000 | 4.000000 | 6.000000 |
| 50% | 8690.0000 | 3.000000 | 1.000000 | 7.000000 | 12.000000 |
| 75% | 13034.5000 | 3.000000 | 1.000000 | 10.000000 | 18.000000 |
| max | 17379.0000 | 4.000000 | 1.000000 | 12.000000 | 23.000000 |

| | holiday | weekday | weathersit | temp | atemp \ |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 17379.000000 | 17379.000000 | 17379.000000 | 17379.000000 | 17379.000000 |
| mean | 0.028770 | 3.003683 | 1.425283 | 0.496987 | 0.475775 |
| std | 0.167165 | 2.005771 | 0.639357 | 0.192556 | 0.171850 |
| min | 0.000000 | 0.000000 | 1.000000 | 0.020000 | 0.000000 |
| 25% | 0.000000 | 1.000000 | 1.000000 | 0.340000 | 0.333300 |
| 50% | 0.000000 | 3.000000 | 1.000000 | 0.500000 | 0.484800 |
| 75% | 0.000000 | 5.000000 | 2.000000 | 0.660000 | 0.621200 |
| max | 1.000000 | 6.000000 | 4.000000 | 1.000000 | 1.000000 |

| | hum | windspeed | casual | registered | count |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 17379.000000 | 17379.000000 | 17379.000000 | 17379.000000 | 17379.000000 |
| mean | 0.627229 | 0.190098 | 35.676218 | 153.786869 | 189.463088 |
| std | 0.192930 | 0.122340 | 49.305030 | 151.357286 | 181.387599 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 1.000000 |
| 25% | 0.480000 | 0.104500 | 4.000000 | 34.000000 | 40.000000 |
| 50% | 0.630000 | 0.194000 | 17.000000 | 115.000000 | 142.000000 |
| 75% | 0.780000 | 0.253700 | 48.000000 | 220.000000 | 281.000000 |
| max | 1.000000 | 0.850700 | 367.000000 | 886.000000 | 977.000000 |

1.1 How many total rentals were there?

```
[85]: df['count'].sum()
```

```
[85]: 3292679
```

1.2 When were the peak hours for rentals?

```
[86]: hourlyCount = df[['hr', 'count']].groupby('hr').sum()
```

```
[87]: hourlyCount.head()
```


```
[87]:
```

| | count |
|----|-------|
| hr | |

```
0    39130
1    24164
2    16352
3     8174
4     4428
```

```
[88]: hourlyCount.plot(xticks=range(0, 25, 2), grid=True, xlim=0, ylim=0)
```

```
[88]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd7573ad490>
```




The peak hours for rentals was at 8 am and 5 pm. This makes sense as 8 am and 5 pm is usually the time for rush hour. These findings suggest that the most common usage of bike rentals is for commuting to work.

1.3 Are there more bike rentals on a weekday vs weekend?

```
[89]: # 0 means Sunday, 1 means Monday, etc.
```

```
df[['weekday', 'count']].groupby('weekday').sum().plot(kind='line', rot=1)
```

```
[89]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd7580a8e90>
```



We can use a hypothesis test to check if there is a significant difference. Technically since the data is taken from 2011 and 2012, it is not a simple random sample of all bike rides for this company.

Null hypothesis: There is no difference in bike rides per day on a weekday vs weekend. (proportion of rides = 0.5)

Alt hypothesis: There are more bike rides per day on a weekday. (proportion of rides > 0.5)

[90]: Create a weekdaybool column and add it to the data set

```
File "<ipython-input-90-4060e1ed4d45>", line 1
Create a weekdaybool column and add it to the data set
^
```

SyntaxError: invalid syntax

```
[92]: counts = df[['weekday', 'count']].groupby('weekday').sum()
counts = counts.reset_index()
```

```
[93]: counts[(counts['weekday'] == 0) | (counts['weekday'] == 6)]
```

```
[93]:   weekday  count
0         0  444027
6         6  477807
```

```
[94]: # output of this function is (z-stat, p-value)
sm.stats.proportions_ztest(2292410, 2292410 + 1000269, 0.5,
↪alternative='two-sided')
```

```
[94]: (774.1944548167168, 0.0)
```

The p-value is 0.0, so there is sufficient evidence to reject the null hypothesis that the proportion of bike rentals on a working day is 0.5. This finding shows that there are more bike rentals on a working day, which aligns with the previous finding that there are more bike rentals around rush hour.

```
[ ]: # single regression of temp vs num of rides
# mult regression predicting # of rides or something
```

1.4 What is the relationship between temperature and number of rides on a day?

```
[24]: df[['dteday', 'temp', 'count']].head()
```

```
[24]:
```

| | dteday | temp | count |
|---|------------|------|-------|
| 0 | 2011-01-01 | 0.24 | 16 |
| 1 | 2011-01-01 | 0.22 | 40 |
| 2 | 2011-01-01 | 0.22 | 32 |
| 3 | 2011-01-01 | 0.24 | 13 |
| 4 | 2011-01-01 | 0.24 | 1 |

According to the data source: “temp: Normalized temperature in Celsius. The values are derived via $(t - tmin)/(tmax - tmin)$, $tmin=-8$, $tmax=+39$ (only in hourly scale)”

The data will be converted to fahrenheit for easier interpretation.

```
[26]: tmin = -8
tmax = 39
```

```
[30]: # temp in celsius, not normalized
celsius = (df['temp'] * (tmax - tmin)) + tmin
```

```
[31]: # temp in fahrenheit, not normalized
fahrenheit = (celsius * 1.8) + 32
fahrenheit
```

```
[31]: 0      37.904
1      36.212
2      36.212
3      37.904
4      37.904
...
```

```
17374    39.596
17375    39.596
17376    39.596
17377    39.596
17378    39.596
Name: temp, Length: 17379, dtype: float64
```

```
[32]: df['temp'] = fahrenheit
```


```
[40]: df[['dteday', 'hr', 'temp', 'count']].head(13)
```

```
[40]:
```

| | dteday | hr | temp | count |
|----|------------|----|--------|-------|
| 0 | 2011-01-01 | 0 | 37.904 | 16 |
| 1 | 2011-01-01 | 1 | 36.212 | 40 |
| 2 | 2011-01-01 | 2 | 36.212 | 32 |
| 3 | 2011-01-01 | 3 | 37.904 | 13 |
| 4 | 2011-01-01 | 4 | 37.904 | 1 |
| 5 | 2011-01-01 | 5 | 37.904 | 1 |
| 6 | 2011-01-01 | 6 | 36.212 | 2 |
| 7 | 2011-01-01 | 7 | 34.520 | 3 |
| 8 | 2011-01-01 | 8 | 37.904 | 8 |
| 9 | 2011-01-01 | 9 | 44.672 | 14 |
| 10 | 2011-01-01 | 10 | 49.748 | 36 |
| 11 | 2011-01-01 | 11 | 48.056 | 56 |
| 12 | 2011-01-01 | 12 | 53.132 | 84 |

```
[42]: df[['temp', 'count']].plot(kind='scatter', x='temp', y='count')
```

```
[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd7519f1f50>
```




We will try to fit the data using linear regression and possibly polynomial regression if a linear model is not a good fit.

```
[44]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[48]: from sklearn.linear_model import LinearRegression
```

```
[54]: plt.figure(figsize=(8, 5))
sns.regplot(x='temp', y='count', data=df)
```

```
[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd754458c10>
```




```
[55]: lm = LinearRegression()  
X = df[['temp']]  
Y = df['count']  
lm.fit(X,Y)  
lm.score(X,Y) #R-squared value
```

```
[55]: 0.16384059523903416
```

```
[57]: plt.figure(figsize=(8, 5))  
sns.residplot(df['temp'], df['count'])
```

```
[57]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd7544a2e10>
```





The linear model has an R-squared value of 0.16 and the residual plot does not show a random scatter making this linear model a poor fit for the data. A 2nd order polynomial regression may be more reasonable because people won't want to ride a bike if the weather is too cold or too hot.

```
[65]: coeff = np.polyfit(df['temp'], df['count'], 2)
      pm = np.poly1d(coeff)
```

```
[69]: plt.plot(df['temp'], df['count'], '.', df['temp'], pm(df['temp']), '-')
```

```
[69]: [<matplotlib.lines.Line2D at 0x7fd75515bdd0>,
      <matplotlib.lines.Line2D at 0x7fd75506fe50>]
```



```
[ ]: # split up for each season
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```